

TRP1 Challenge Week 1 — Interim Report

1. How the VS Code Extension Works

1.1 High-Level Architecture

Roo Code is a VS Code extension that provides an AI-powered coding assistant running entirely inside the editor. Its architecture follows a strict **three-layer separation**:

Layer	Responsibility	Key Files
Webview UI	Presentation layer — renders chat, diff views, settings. Communicates with the extension host via <code>postMessage</code> . Cannot access the file system or network directly.	<code>webview-ui/src/</code> (React + Vite)
Extension Host	Core logic — receives messages from the webview, manages tasks, invokes LLM APIs, executes tools, handles approval gates, and manages MCP servers.	<code>src/extension.ts</code> , <code>src/core/</code>
LLM Providers	External API layer — sends system prompt + conversation history to model providers (Anthropic, OpenAI, etc.) and streams back assistant responses.	<code>src/api/providers/</code>

1.2 Extension Activation

When VS Code activates the extension (`src/extension.ts:120`):

1. An output channel is created for logging.
2. Telemetry, i18n, terminal registry, and network proxy are initialized.
3. A `ClineProvider` is instantiated — this is the central controller that bridges the webview and the task runtime.
4. The provider is registered as a `WebViewProvider` on VS Code's sidebar.
5. Commands, code actions, and terminal actions are registered.
6. Cloud services and MCP server manager are initialized.

1.3 Message Flow: User Prompt → Tool Execution

User types prompt in Webview

```
|  
v  
postMessage("newTask", text)  
|
```

```

    v
ClineProvider.setWebviewMessageListener()  [ClineProvider.ts:1326]
|
v
webviewMessageHandler()  [webviewMessageHandler.ts:89]
|  case "newTask" -> provider.createTask(text)
|
v
Task created  [src/core/task/Task.ts]
|  Task.start() -> recursivelyMakeClineRequests()
|
v
System prompt built  [src/core/prompts/system.ts:41]
Tools built  [src/core/task/build-tools.ts:82]
|
v
LLM API call  [src/api/providers/*.ts]
|  Streams assistant response
|
v
presentAssistantMessage()  [presentAssistantMessage.ts:61]
|  Processes content blocks sequentially
|  For each tool_use block:
|    1. validateToolUse() -- mode/permission check
|    2. askApproval() -- human-in-the-loop gate
|    3. tool.execute() -- runs the tool
|    4. pushToolResult() -- returns result to conversation
|
v
Next LLM turn or completion

```

1.4 Tool System

Tools are the agent's interface to the environment. Each tool is:

- **Defined** as an OpenAI-compatible function schema in `src/core/prompts/tools/native-tools/`.
- **Filtered** per mode (architect, code, ask, etc.) in `src/core/prompts/tools/filter-tools-for-mode.ts`.
- **Dispatched** in the tool switch block of `presentAssistantMessage.ts` (~line 678).
- **Implemented** as a class extending `BaseTool` in `src/core/tools/`.

Current native tools include: `write_to_file`, `execute_command`, `read_file`, `list_files`, `search_files`, `apply_diff`, `apply_patch`, `edit_file`, `search_replace`, `codebase_search`, `use_mcp_tool`, `access_mcp_resource`, `ask_followup_question`, `attempt_completion`, `switch_mode`, `new_task`, `update_todo_list`, `run_slash_command`, `skill`, `generate_image`.

1.5 Approval and Auto-Approval

Before any tool executes, the extension runs an approval gate:

- **Manual approval:** The user sees a prompt (diff view for writes, command preview for execution) and clicks Approve/Reject.
- **Auto-approval:** Configurable policies in `src/core/auto-approval/index.ts:47` can auto-approve safe operations (reads) while requiring manual approval for destructive ones (writes, commands).

- **Protected files:** `rooIgnoreController` and `rooProtectedController` add file-level access restrictions.
-

2. Code and Design Architecture of the Agent (Phase 0 — ARCHITECTURE_NOTES.md)

2.1 End-to-End Execution Path

The following traces the exact code path for every tool call:

Step	What Happens	Exact Location
1	System prompt generated	<code>src/core/prompts/system.ts:41</code> <code>(generatePrompt)</code>
2	Tool definitions built	<code>src/core/prompts/tools/native-tools/index.ts:10</code> <code>(getNativeTools)</code>
3	Per-request tool filtering	<code>src/core/task/build-tools.ts:82</code> <code>(buildNativeToolsArrayWithRestrictions)</code>
4	Assistant output processed	<code>src/core/assistant-message/presentAssistantMessage.ts:597</code>
5	Tool permission check	<code>→ validateToolUse()</code>
6	Human approval gate	<code>presentAssistantMessage.ts:494</code> <code>(askApproval callback)</code>
7	Tool result callback	<code>presentAssistantMessage.ts:449</code> <code>(pushToolResult callback)</code>
8	Tool dispatch switch	<code>presentAssistantMessage.ts:678</code>
9	<code>write_to_file</code> dispatch	<code>presentAssistantMessage.ts:679</code> <code>→</code> <code>WriteToFileTool.execute()</code>
10	<code>execute_command</code> dispatch	<code>presentAssistantMessage.ts:764</code> <code>→</code> <code>ExecuteCommandTool.execute()</code>

2.2 Detailed Sequence: `write_to_file`

1. Assistant emits `tool_use(write_to_file, {path, content})`.
2. `validateToolUse` checks mode permissions.
3. `askApproval` presents a diff view to the user.
4. `WriteToFileTool.execute()` runs at `src/core/tools/WriteToFileTool.ts:29`.
5. Access checks: `rooIgnoreController.validateAccess()` (line 50), `rooProtectedController.isWriteProtected()` (line 58).
6. Diff is computed, shown in editor, and user approves.
7. File is saved to disk.
8. Result is pushed through `pushToolResult` back into the conversation.

2.3 Detailed Sequence: execute_command

1. Assistant emits `tool_use(execute_command, {command, cwd})`.
2. `validateToolUse` checks mode permissions.
3. `askApproval` shows the command for approval.
4. `ExecuteCommandTool.execute()` runs at `src/core/tools/ExecuteCommandTool.ts:34`.
5. Command is validated against `rootIgnoreController`.
6. Terminal execution with timeout and allowlist checks.
7. Output is streamed to webview and final result becomes `tool_result`.

2.4 Hook Injection Points Identified

Injection Point	Location	Purpose
Pre-hook	Before tool dispatch switch at <code>presentAssistantMessage.ts:678</code>	Enforce intent handshake, scope validation, command classification
Post-hook	After tool execution at <code>presentAssistantMessage.ts:920</code>	Record trace entries, compute content hashes, update sidecar files
Prompt protocol	<code>src/core/prompts/system.ts:41</code>	Add handshake instructions requiring <code>select_active_intent</code> before mutations
Tool contract	<code>src/core/prompts/tools/nativeTools/index.ts:42</code>	Registers tool <code>select_active_intent</code> tool definition

2.5 Current Gaps Against Target Architecture

- No `select_active_intent` tool exists yet.
- No `.orchestration/` sidecar integration in the tool execution path.
- No intent-linked `agent_trace.jsonl` write path.
- No explicit mutation classification (`AST_REFATOR`, `INTENT_EVOLUTION`).
- No stale-write optimistic lock guard in the write path.

3. Architectural Decisions for the Hook

Decision 1: Middleware Boundary at Tool Execution Loop

- **Decision:** Place hook execution at the extension host tool loop — before and after tool dispatch.
- **Why:** `presentAssistantMessage.ts` is the single central path where every tool call passes. Hooking here guarantees 100% interception coverage.
- **Location:** `src/core/assistant-message/presentAssistantMessage.ts`.
- **Tradeoff:** Tight coupling to current execution loop shape; requires careful refactoring if the loop changes.

Decision 2: Two-Stage Handshake for Mutating Tools

- **Decision:** Mutating tools (`write_to_file`, `execute_command`, `apply_diff`, etc.) must require an active intent selected through `select_active_intent(intent_id)` before execution.
- **Why:** Enforces intent-code traceability. Every code mutation has a traceable business reason.
- **Tradeoff:** Adds one extra tool call before writes/commands. Read-only tools remain unblocked.

Decision 3: Fail-Closed for Governance, Fail-Safe for Observability

- **Decision:** Pre-hooks block on policy violations (fail-closed). Post-hooks do not block user progress if telemetry/trace writes fail (fail-safe).
- **Why:** Policy enforcement must be strict — an unaudited mutation cannot be allowed. But logging failures should not corrupt the main task loop.
- **Tradeoff:** Post-hook failures must be surfaced clearly for audit (logged as errors, not silently dropped).

Decision 4: Sidecar Data Model in `.orchestration/`

- **Decision:** Persist intent and trace artifacts in workspace-local sidecar files under `.orchestration/`.
- **Why:** Keeps audit metadata close to code changes. Files can be committed to version control alongside the code they describe.
- **Tradeoff:** File-based writes need append discipline and eventual compaction for the JSONL ledger.

Decision 5: Append-Only Trace Ledger

- **Decision:** `agent_trace.jsonl` uses append-only records with content hashes.
- **Why:** Prevents silent history rewrites and supports deterministic auditing. Content hashes ensure spatial independence (if code lines move, the hash still validates).
- **Tradeoff:** Requires periodic compaction strategy if the file grows large.

Decision 6: Minimal Clean Hook Module

- **Decision:** Keep hook contracts and engine isolated in `src/hooks/`, separate from the core execution loop.
 - **Why:** Prevents spaghetti logic in the main tool dispatch. Enables composable, testable hooks that can be registered independently.
 - **Tradeoff:** Requires explicit wiring points at integration time.
-

4. Diagrams and Schemas of the Hook System

4.1 Runtime Boundary Diagram

```
graph LR
    U["User in VS Code"] --> WV["Webview UI"]
    WV -->|postMessage| EH["Extension Host"]
```

```

EH -->|assistant output| LOOP[presentAssistantMessage loop]
LOOP --> PRE[Pre-Tool Hooks]
PRE --> DISPATCH[Tool Dispatch]
DISPATCH --> TOOL[Native Tool Execute]
TOOL --> POST[Post-Tool Hooks]
POST --> ORCH[.orchestration sidecar]
POST --> RESP[tool_result to UI]
RESP --> WV

```

4.2 Two-Stage Handshake Sequence

```

sequenceDiagram
    participant User
    participant Agent
    participant HookEngine
    participant Sidecar as .orchestration

    User->>Agent: "Refactor auth middleware"
    Agent->>HookEngine: tool_use(select_active_intent, intent_id)
    HookEngine-->>Sidecar: read active_intents.yaml
    Sidecar-->>HookEngine: intent constraints + scope
    HookEngine-->>Agent: inject context patch, allow
    Agent->>HookEngine: tool_use(write_to_file, ...)
    HookEngine-->>HookEngine: pre-hook policy checks
    HookEngine-->>Agent: allow or block
    Agent->>HookEngine: tool_result
    HookEngine-->>Sidecar: append agent_trace.jsonl

```

4.3 Hook Engine Type Schema

```

// src/hooks/types.ts

export type HookPhase = "preToolUse" | "postToolUse"

export interface HookContext {
    taskId: string
    toolName: string
    toolArgs: Record<string, unknown>
    cwd?: string
    timestamp: string
}

export type BlockCode =
    | "INTENT_REQUIRED"
    | "SCOPE_VIOLATION"
    | "DESTRUCTIVE_BLOCKED"
    | "HOOK_ERROR"

```

```

export type HookDecision =
  | { allow: true; contextPatch?: Record<string, unknown> }
  | { allow: false; reason: string; code: BlockCode }

export interface PreToolHook {
  name: string
  run(context: HookContext): Promise<HookDecision> | HookDecision
}

export interface PostToolHookContext extends HookContext {
  toolResult: unknown
  changedFiles: string[]
}

export interface PostToolHook {
  name: string
  run(context: PostToolHookContext): Promise<void> | void
}

```

4.4 HookEngine Implementation

```

// src/hooks/HookEngine.ts

export class HookEngine {
  private readonly preHooks: PreToolHook[] = []
  private readonly postHooks: PostToolHook[] = []

  registerPreHook(hook: PreToolHook): void { ... }
  registerPostHook(hook: PostToolHook): void { ... }

  async runPreHooks(context: HookContext): Promise<HookDecision> {
    // Runs all pre-hooks sequentially.
    // If any hook returns { allow: false }, execution is blocked immediately.
    // Context patches from passing hooks are merged.
    // Errors in hooks are caught and returned as HOOK_ERROR blocks.
  }

  async runPostHooks(context: PostToolHookContext): Promise<{ errors: string[] }> {
    // Runs all post-hooks sequentially.
    // Errors are collected but do not block the main flow (fail-safe).
  }
}

```

4.5 active_intents.yaml Schema

```

# .orchestration/active_intents.yaml
active_intents:
  - id: "INT-001"
    name: "JWT Authentication Migration"
    status: "IN_PROGRESS"
    owned_scope:
      - "src/auth/**"
      - "src/middleware/jwt.ts"
    constraints:
      - "Must not use external auth providers"
      - "Must maintain backward compatibility with Basic Auth"
    acceptance_criteria:
      - "Unit tests in tests/auth/ pass"

```

4.6 agent_trace.jsonl Record Schema

```
{
  "id": "uuid-v4",
  "timestamp": "2026-02-18T12:00:00Z",
  "intent_id": "INT-001",
  "mutation_class": "AST_REFACTOR",
  "vcs": { "revision_id": "git_sha_hash" },
  "files": [
    {
      "relative_path": "src/auth/middleware.ts",
      "conversations": [
        {
          "url": "session_log_id",
          "contributor": {
            "entity_type": "AI",
            "model_identifier": "claude-3-5-sonnet"
          },
          "ranges": [
            {
              "start_line": 15,
              "end_line": 45,
              "content_hash": "sha256:a8f5f167f44f4964e6c998dee827110c"
            }
          ],
          "related": [
            {
              "type": "specification",
              "value": "INT-001"
            }
          ]
        }
      ]
    }
  ]
}
```

```

        ]
    }
}

```

5. Repository Deliverable Status

5.1 `src/hooks/` Directory Contents

File	Purpose
<code>types.ts</code>	Shared contracts — <code>PreToolHook</code> , <code>PostToolHook</code> , <code>HookDecision</code> , <code>BlockCode</code>
<code>HookEngine.ts</code>	Composable middleware runner with fail-closed pre-hooks and fail-safe post-hooks
<code>builtin/RequireIntentPreHook.ts</code>	Pre-hook: blocks mutating tools unless <code>intent_id</code> is provided
<code>builtin/TraceMutationPostHook.ts</code>	Post-hook: placeholder for mutation trace recording to <code>agent_trace.jsonl</code>
<code>index.ts</code>	Public exports for the hook module
<code>README.md</code>	Integration guide for wiring hooks into the runtime

5.2 Documentation Files

File	Content
<code>ARCHITECTURE_NOTES.md</code>	Phase 0 archaeological dig: exact file:line mappings for tool loop, prompt builder, approval flow, and injection points
<code>HOOK_ARCHITECTURAL_DECISIONS.md</code>	6 explicit design decisions with rationale and tradeoffs
<code>HOOK_DIAGRAMS_AND_SCHEMAS.md</code>	Mermaid diagrams (runtime boundary, handshake sequence) and data schemas
<code>INTERIM_REPORT.md</code>	This report
<code>INTERIM_REPORT.pdf</code>	PDF export of this report

5.3 `.orchestration/` Directory

File	Status
<code>active_intents.yaml</code>	Created (to be populated with example intents)
<code>agent_trace.jsonl</code>	Planned for Phase 3 implementation
<code>intent_map.md</code>	Planned for Phase 3 implementation

6. Interim Scope Note

This interim submission delivers the **Phase 0 (Archaeological Dig)** deliverables and the **Phase 1/2 scaffold**:

- Complete architecture mapping with exact insertion points for hook integration.
- Clean, isolated `src/hooks/` module with typed contracts, engine, and starter hooks.
- Documented architectural decisions and visual schemas.

Full runtime integration (wiring hooks into `presentAssistantMessage.ts`), the `select_active_intent` tool, trace persistence, scope enforcement, and parallel orchestration are planned for the final submission phases.